



Bruce Marron <marron.bruce.mx@gmail.com>

Important Update: DeepSeek API Upgraded to DeepSeek-V3.1 & Pricing Adjustment

1 message

DeepSeek <support@sc.mail.deepseek.com>
Reply-To: DeepSeek <support@sc.mail.deepseek.com>
To: marron.bruce.mx@gmail.com

Thu, Aug 21, 2025 at 1:08 AM



Dear Developers,

We are pleased to announce that DeepSeek API has been fully upgraded to our new premier model: DeepSeek-V3.1. Key enhancements include:

1. **Unified Dual-Mode Support**

DeepSeek-V3.1 combines reasoning (deepseek-reasoner) and non-reasoning (deepseek-chat) capabilities in a single model. Legacy models will be retired.

2. **Performance Upgrades**

Significant improvements in Agent tasks, reasoning efficiency, and 128K-token context handling.

3. **The Strict-mode for Function Calling**

Beta API now supports strict-mode for function calling, ensuring the output always complies with the defined JSON schema.

4. **Anthropic API Compatibility**

The API is now compatible with Anthropic API, allowing seamless integration with Claude Code.

Pricing Adjustment (Effective Sept. 4, 2025, 16:00 UTC):

- Input tokens: \$0.07/million (cache hit) / \$0.56/million (cache miss)
- Output tokens: \$1.12/million
- Off-peak discounts discontinued simultaneously

Current pricing remains until the effective date. Plan usage accordingly.

For more details, please visit [DeepSeek API Documentation](#).

Thank you for your continued trust in DeepSeek.

Best regards,

The DeepSeek Team

August 20, 2025

点击这里取消订阅
click to unsubscribe